

# Programming Assignment 3 Part B

Binghan Geng

A20482350

[bgeng1@hawk.iit.edu](mailto:bgeng1@hawk.iit.edu)

Yu Li

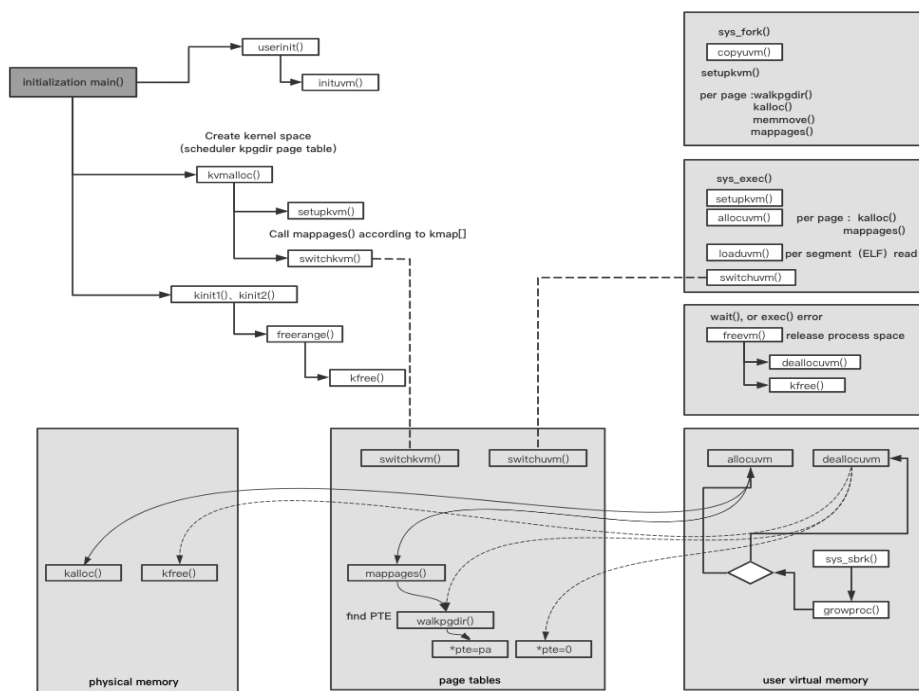
A20496405

[yli385@hawk.iit.edu](mailto:yli385@hawk.iit.edu)

## 1. System calls to share a memory page

### 1.1 Overview : xv6 memory management

XV6 memory management is divided into two stages: initialization and runtime. It involves physical page management, virtual memory space management, virtual and real mapping page tables. On top of these three operations, xv6 implements memory management initialization operations and processing dynamic runtime operations. The following figure is the whole process of memory management organized according to xv6, and the call link of the code.



## 1.2 Implementation of xv6 memory sharing

### 1.2.1 The basic program structure

The simplified scheme of process mapping shared memory page is as shown in the figure below, process A, B, C realize data reading and writing through the mapping relationship between virtual address and physical address.

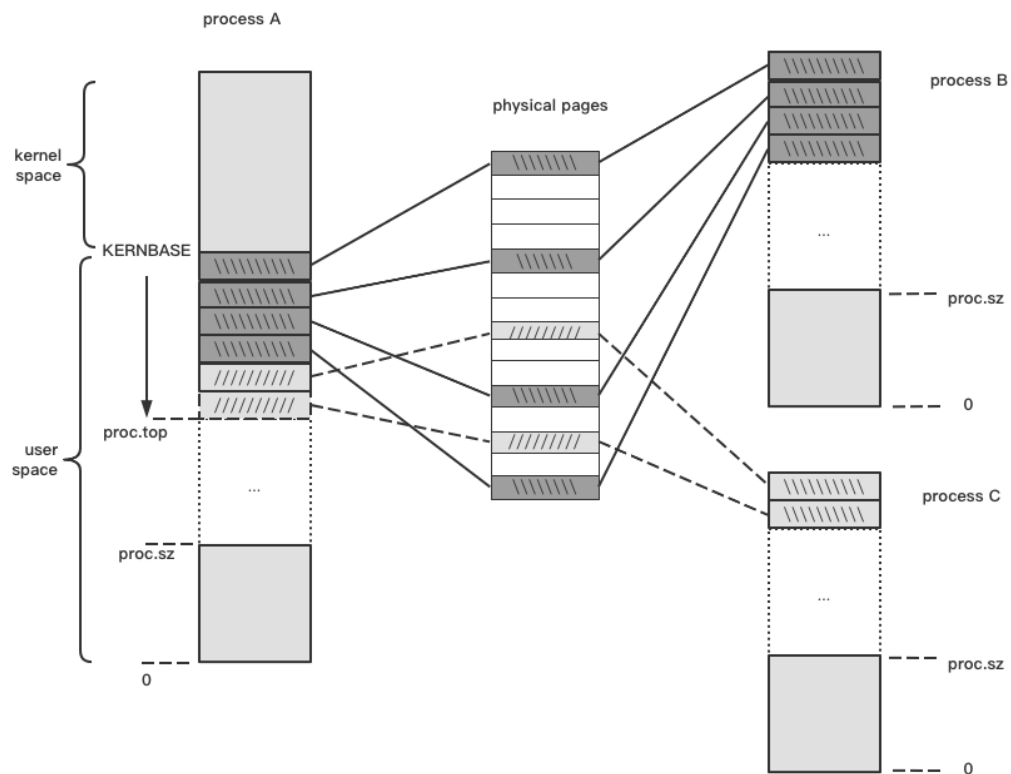


Figure: Program shared memory diagram

### 1.2.2 The basic data structure

The program have **A maximum of shared memory areas**, using array of record structure,each member is a “sharemempage” structure, the members inside include the reference count (the number of times the process is mapped), the physical page mapped in this interval (code 1-1. line 390-394).

**A maximum of physical pages corresponding to the key**, using array of integer “page\_nums[NKEY]” , cooperate with sharemempage(code 1-1. line 398).

**A Lock “shmlock”** , using spinlock struct for exclusive alloc/dealloc memory(code 1-1. line 396).

```

389
390 struct sharemempage
391 {
392     int refcount;           //Reference count for each key
393     void* physaddr[MAX_SHM_PGNUM]; //Corresponding to the physical address of per page
394 };
395
396 struct spinlock shmlock;           //Locks for exclusive access
397 struct sharemempage shmarray[NKEY]; //Array of Share memory page record
398 int page_nums[NKEY]; // Number of physical pages corresponding to the key
399

```

code 1-1 vm.c

The program also defines 3 variables in proc structure (proc.h): **top**, is used to mark the upper limit of the user space of the process, originally KERNELBASE, but the shared content is allocated at the top of the user space, recording this upper limit can be used normally sbrk(). **shm\_key[NKEY]**, is a array, which shared memory is used at which position. **shm\_va[NKEY]**, virtual address space of used pages.

```

50
51 // Per-process state
52 struct proc {
53     uint sz;           // Size of process memory (bytes)
54     pde_t* pgdir;      // Page table
55     char *kstack;      // Bottom of kernel stack for this process
56     enum procstate state; // Process state
57     int pid;           // Process ID
58     struct proc *parent; // Parent process
59     struct trapframe *tf; // Trap frame for current syscall
60     struct context *context; // switch() here to run process
61     void *chan;         // If non-zero, sleeping on chan
62     int killed;         // If non-zero, have been killed
63     struct file *ofile[NOFILE]; // Open files
64     struct inode *cwd;  // Current directory
65     char name[16];      // Process name (debugging)
66
67     uint top;           //Current top of VA (the high end of the process's address space)
68     uint shm_keys[NKEY]; //Keys that the process has called GetSharedPage
69     void* shm_va[NKEY]; //Virtual address space of used pages
70
71 };
72

```

code 1-2 proc.h

### 1.2.3 Implementation details

**Setp1:** Add a new system call void \* GetSharedPage(key, num\_pages).

The parameter **key**, is used to specify which of the shared memory areas to obtain, **num\_pages** represents the number of the shared memory pages. This system call returns the virtual address corresponding to the shared memory pages.

When the process calls GetSharedPage(key, num\_pages), it is divided into two cases for processing according to whether the shared memory already exists:

1. If the shared memory area corresponding to `shmall[key]` has not been created yet, initialize `shmall[key]` according to the size, allocate the corresponding physical memory, establish a page table mapping, and return the virtual address of the memory.

2. If the `shmall[key]` shared memory area already exists, the new page table entry `pte` is bound to the physical address of `shmall[key]` pages, and the corresponding virtual address is returned. At this time, the incoming size parameter is ignored.

Similarly, we also need to use a series of functions such as the initialization operation function `shminit()` of the shared memory area. the newly added function is now in the `vm.c` file and need to be declared in `defs.h` so that other codes can call them. As shown in code 2-1.

```
178 void      shminit();
179 void*      GetSharedPage(uint key, uint num_pages);
180 int        FreeSharedPage(uint key);
181
```

code 2-1 defs.h

## Setp2: Initialize

Add a `shminit` method in `vm.c` to initialize the shared variable array `shmall[]`.

```
464 //initialize the shared variable array shmall[].
465 void
466 shminit()
467 {
468     initlock(&shmlock, "shmplock"); //Initializing locks
469
470     for (int i = 0; i < NKEY; i++) //Initialize shmall
471     {
472         shmall[i].refcount = 0;
473         page_nums[i] = 0;
474         for (j = 0; j < MAX_SHM_PGNUM; j++) {
475             physaddr[j] = 0;
476         }
477     }
478 }
479
```

code 2-2 vm.c

Then call the initialization function in the main() method in main.c, so that the preparation of shared variables can be completed when the system starts.

```
37 startothers(); // start other processors
38 kinit2(P2V(4*1024*1024), P2V(PHYSTOP)); // must come after startothers()
39 shminit(); // initialize the shared memory page
40 userinit(); // first user process
41 mpmain(); // finish this processor's setup
42 }
```

code 2-3 main.c

```
35 static struct proc*
36 allocproc(void)
37 {
38     struct proc *p;
39     char *sp;
40
41     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
42         if(p->state == UNUSED)
43             goto found;
44     return 0;
45
46 found:
47     p->state = EMBRYO;
48     p->pid = nextpid++;
49     p->top = KERNBASE; //Initialize shm, at first shm overlaps with KERNBASE
50
51
52     for (int i = 0; i < NKEY; i++) {
53         p->shm_keys[i] = 0;
54     }
```

code 2-4 proc.c

### Setp3: Map shared memory to process space

Next, need to implement GetSharedPage(key, num\_pages) that maps the shared memory pages area to the process space. As mentioned earlier (NKEY = 8), the parameter key is a subscript (0-7) of the shared memory, and num\_pages is the number of pages to be allocated, and returns the address of the shared memory. The implementation needs to be divided into three situations to implement:

1. If the shared memory area has been mapped in the process space, the address is returned directly.
2. If the shared memory area has not been created, the reference count of the system shmarray[key] is still 0, we need to create the memory area (including allocating pages and establishing page table mapping).

The method `allocshmex()`, the working principle of `allocuvm()` is the same, the difference is that our memory usage is allocated from high addresses to low addresses, and method `deallocshmex()` is also the same.

```

439 // This method is basically the same as the allocuvm implementation, the difference is
440 // that our memory usage is allocated from high addresses to low addresses.
441 int
442 allocshmex(pde_t *pgdir, uint oldva, uint newva, uint sz, void *phyaddr[MAX_SHM_PGNUM])
443 {
444     char *mem;
445     uint a;
446
447     if(oldva > KERNBASE || newva < sz)
448         return 0;
449     a = newva;
450     for (int i = 0; a < oldva; a+=PGSIZE, i++)
451     {
452         mem = kalloc(); //Assigning physical page frames
453         if(mem == 0){
454             // cprintf("allocshm out of memory\n");
455             deallocshmex(pgdir, newva, oldva);
456             return 0;
457         }
458         memset(mem, 0, PGSIZE);
459         mappages(pgdir, (char*)a, PGSIZE, pa: (uint)V2P(mem), perm: PTE_W|PTE_U); //页表映射
460         phyaddr[i] = (void *)V2P(mem);
461         // cprintf("allocshm : %x\n", a);
462     }
463     return newva;
464 }

```

code 2-5 vm.c

```

415 // This method is basically the same as the deallocuvm implementation, the difference is
416 // that our memory usage is deallocated from low addresses to high addresses.
417 int
418 deallocshmex(pde_t *pgdir, uint oldva, uint newva)
419 {
420     pte_t *pte;
421     uint a, pa;
422     if(newva <= oldva)
423         return oldva;
424     a = (uint)PGROUNDDOWN(a: newva - PGSIZE); //
425     for (; oldva <= a; a-=PGSIZE)
426     {
427         pte = walkpgdir(pgdir, (char*)a, alloc: 0);
428         if(pte && (*pte & PTE_P) != 0){
429             pa = PTE_ADDR(pte: *pte);
430             if(pa == 0){
431                 panic("kfree");
432             }
433             *pte = 0;
434         }
435     }
436     return newva;
437 }

```

code 2-6 vm.c

3. If the system already has corresponding shared memory, just map it to the process space.

```
502 void *
503 GetSharedPage(uint key, uint num_pages) {
504     void *phyaddr[MAX_SHM_PGNUM];
505     uint addr = 0;
506     if (key < 0 || NKEY <= key || num_pages < 0 || MAX_SHM_PGNUM < num_pages) //Calibration parameters
507         return (void *) -1;
508     addr = proc->top;
509     // If the current process has already mapped the shared memory of the key, return the address directly.
510     // Do not support multiple mappings of the same shared memory area in the same process.
511     if (proc->shm_keys[key] == 1) {
512         return proc->shm_va[key];
513     }
514     acquire(&shmlock);
515     // If the system has not yet created the shared memory corresponding to this key, then allocate the memory and map.
516     if (shmarray[key].refcount == 0) {
517         addr = allocshmex(proc->pgdir, addr, newva: addr - num_pages * PGSIZE, proc->sz, phyaddr);
518         //The new allocshmex() allocates memory and maps it, the same principle as allocvmm()
519         if (addr == 0) {
520             release(&shmlock);
521             return (void *) -1;
522         }
523         proc->shm_va[key] = (void *) addr;
524         shmadd(key, num_pages, phyaddr); //Fill the new memory area information into the shmarray[8] array
525     } else {
526         //If the key is not held and the shared memory corresponding to this key is already allocated in the system,
527         // then it is mapped directly.
528         for (int i = 0; i < num_pages; i++) {
529             phyaddr[i] = shmarray[key].physaddr[i];
530         }
531         num_pages = page_nums[key];
532         //The mapshmex method creates a new mapping
533         if ((addr = mapshmex(proc->pgdir, addr, newva: addr - num_pages * PGSIZE, proc->sz, phyaddr)) == 0) {
534             release(&shmlock);
535             return (void *) -1;
536         }
537         proc->shm_va[key] = (void *) addr;
538         shmarray[key].refcount++; //Reference Count +1
539     }
540     proc->top = addr;
541     proc->shm_keys[key] = 1;
542     release(&shmlock);
543     return (void *) addr;
544 }
```

GetSharedPage

code 2-7 vm.c

**Setp4: Remove/Unmap the shared memory area:** int FreeSharedPage(uint key).

FreeSharedPage(key) is used to release / unmap the shared memory pages, and is also called to release the physical memory to completely destroy the Shared memory. If the refcount of a shared memory area becomes 0, it must be called to release the physical memory to completely destroy the shared memory. And because the key addition is disorderly, if you follow the virtual address released by the key, the calculation is very complicated. So when the memory map of the current process (shm\_keys[]) is all zero, all are released at one time.

```

549 FreeSharedPage(uint key) {
550     if (key < 0 || NKEY <= key) {
551         cprintf("key error\n");
552         return -1;
553     }
554     if (proc->shm_keys[key] != 1) {
555         return -1;
556     }
557     acquire(&shmlock);
558     proc->shm_keys[key] = 0;
559
560     cprintf("FreeSharedPage: key is %d\n", key);
561     struct sharemempage *shmem = &shmmarray[key];
562     cprintf("FreeSharedPage: refcount is %d\n", shmem->refcount);
563     shmem->refcount--;
564     cprintf("FreeSharedPage: page_nums is %d\n", page_nums[key]);
565     if (shmem->refcount == 0) {
566         for (int i = 0; i < page_nums[key]; i++) {
567             kfree((char *) P2V(shmem->physaddr[i])); //Recycle physical memory, page by page, frame by frame
568         }
569     }
570     //If the current key without proc are 0, then free the space
571     int flag = 0;
572     for (int i = 0; i < NKEY; i++) {
573         if (proc->shm_keys[i] == 1) {
574             flag++;
575             break;
576         }
577     }
578     if (flag == 0) {
579         // Free the user space memory, because the add is unordered, the calculation of the free virtual
580         // address is very complicated, only a one-time free
581         dealloclshmem(proc->pgdir, proc->top, KERNBASE);
582         cprintf("Free the user space memory.\n");
583     }
584     release(&shmlock);
585     return 0;
586 }

```

code 2-8 vm.c